

# [320] Welcome + First Lecture

## [reproducibility]

Tyler Caraza-Harter

# Introductions

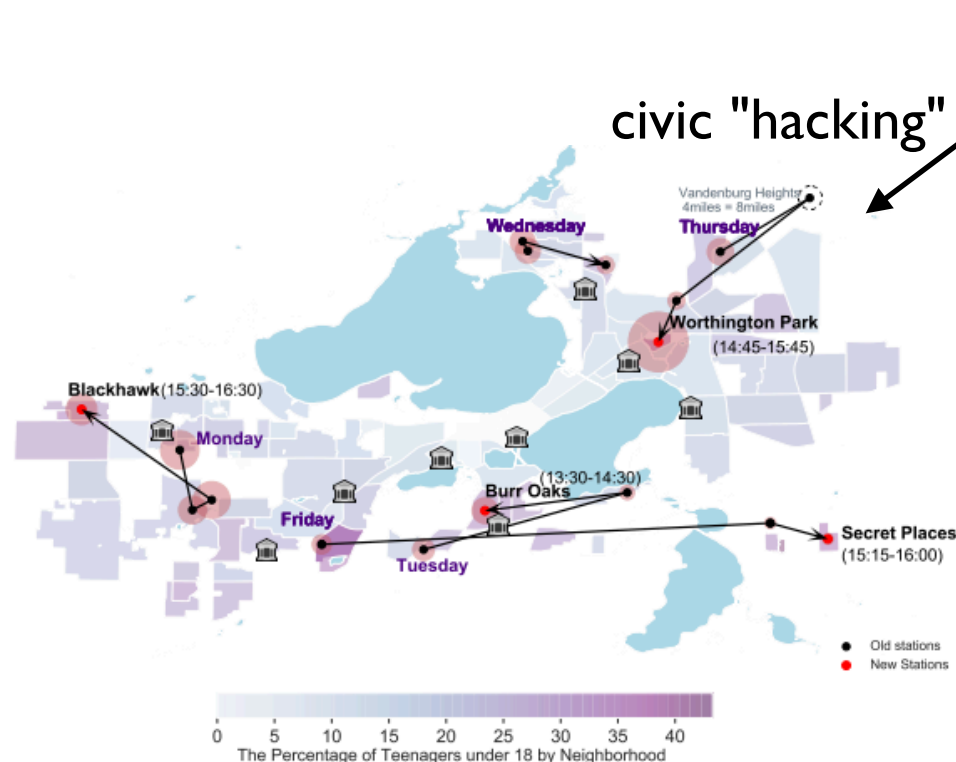
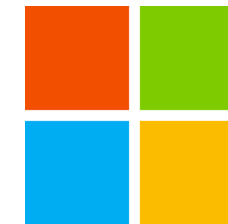
Tyler Caraza-Harter

- Long time Badger
- Email: [tharter@wisc.edu](mailto:tharter@wisc.edu)
- Just call me “Tyler” (he/him)



Industry experience

- Worked at Microsoft on SQL Server and Cloud
- Other internships/collaborations: Qualcomm, Google, Facebook, Tintri



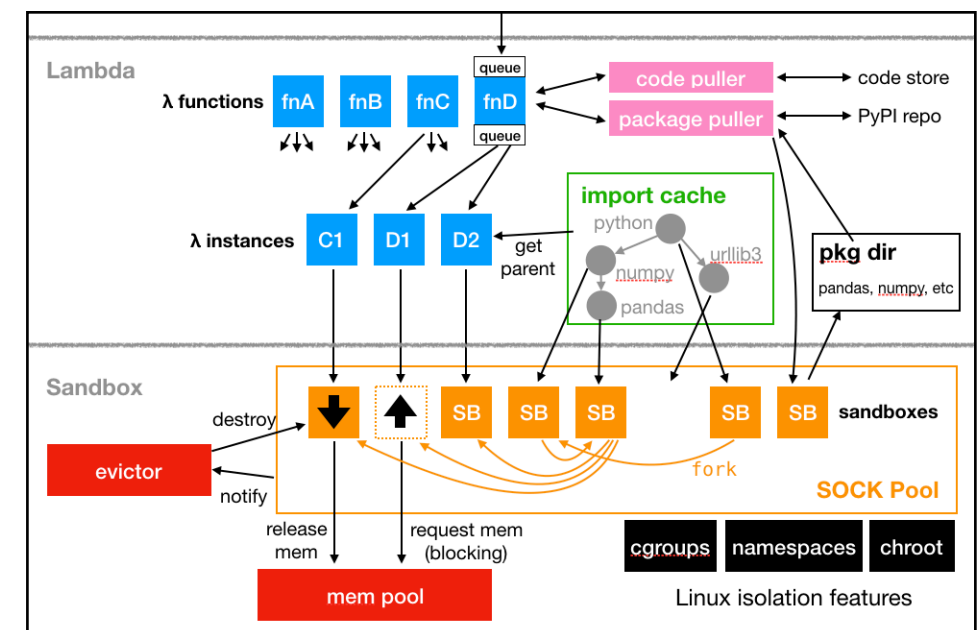
Plot by [Zishan Bai & Dingyi Zhou](#) (previous students)

More: <https://wisc-ds-projects.github.io>

*interests*

civic "hacking"

OpenLambda



# Who are You?

Year in school?

- 1st year? 2nd? Junior/senior? Grad student?

Area of study

- Natural science, social science, engineering, business, statistics, data science, other?

What CS courses have people taken before?

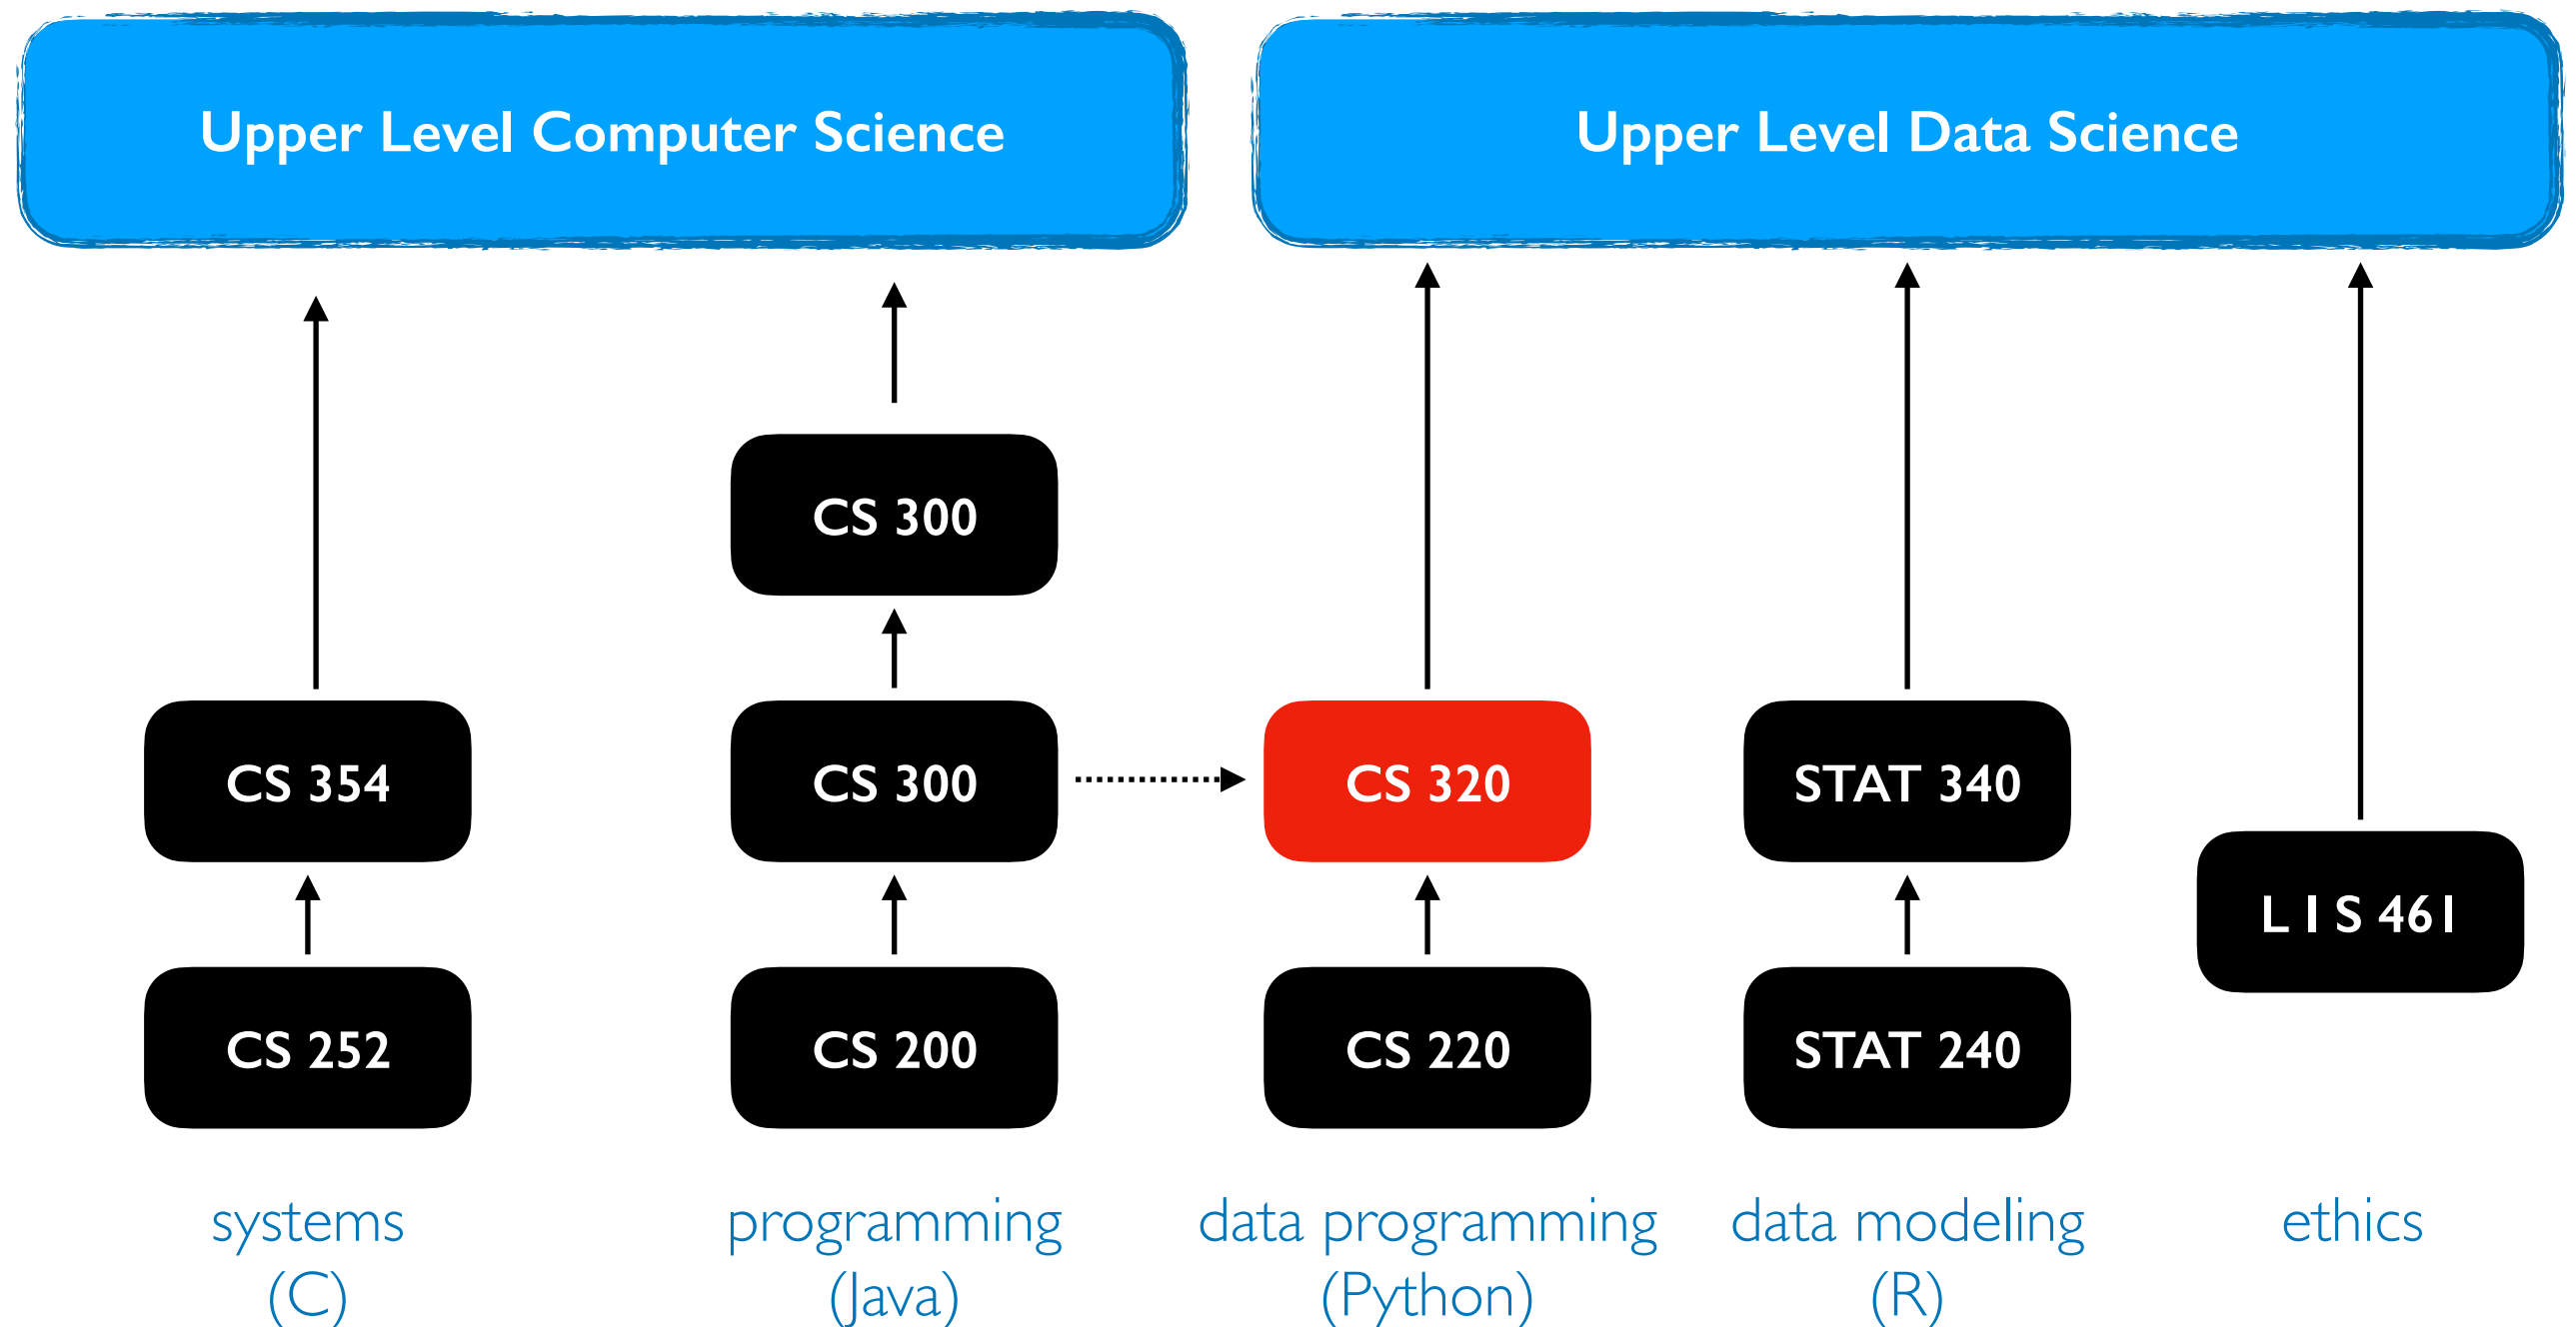
- CS 220/301? CS 200? CS 300? CS 354?

Please fill this form (**due today**): <https://forms.gle/pkchjN6UJq8oVhe68>.

Why?

- Help me get to know you
- Get participation credit
- Group formation

# Related courses



PI (Project I) and other resources will help 300-to-320 students.

# Welcome to Data Programming II, in person!

Builds on CS 30+ 220. <https://stat.wisc.edu/undergraduate-data-science-studies/>

## CS 220

getting results  
writing correct code  
using objects  
functions: `f(obj)`  
lists+dicts  
analyzing datasets  
plots  
tabular analysis

## CS 320

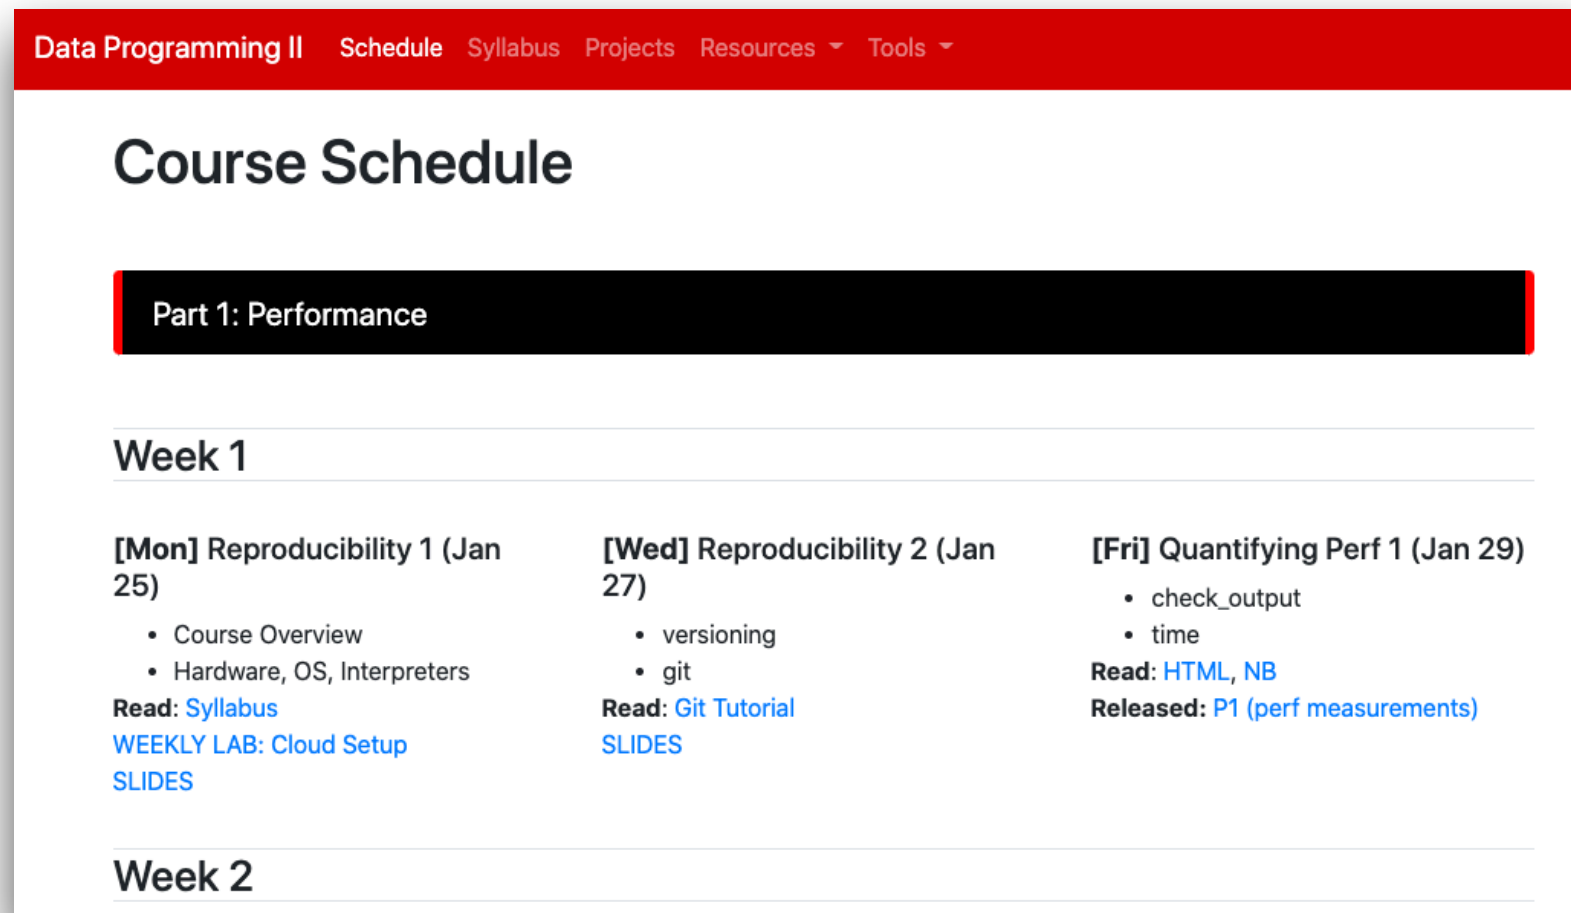
getting **reproducible** results  
writing **efficient** code  
designing **new types** of objects  
**methods**: `obj.f()`  
graphs+trees  
**collecting**+analyzing datasets  
animated visualizations  
**simple machine learning**



# Course Logistics

# Course Website

It's here: <https://tyler.caraza-harter.com/cs320/s22/schedule.html>



The screenshot shows a web page titled "Course Schedule" for "Data Programming II". The navigation bar includes links for "Schedule", "Syllabus", "Projects", "Resources", and "Tools". A black bar highlights "Part 1: Performance". Under "Week 1", there are three columns for different days: Monday, Wednesday, and Friday. Each column lists topics, a "Read" link, and a "Released" link. Week 2 is partially visible at the bottom.

Week 1	Week 2
<b>[Mon] Reproducibility 1 (Jan 25)</b> <ul style="list-style-type: none"><li>• Course Overview</li><li>• Hardware, OS, Interpreters</li></ul> <b>Read:</b> <a href="#">Syllabus</a> <b>WEEKLY LAB:</b> <a href="#">Cloud Setup</a> <b>SLIDES</b>	
<b>[Wed] Reproducibility 2 (Jan 27)</b> <ul style="list-style-type: none"><li>• versioning</li><li>• git</li></ul> <b>Read:</b> <a href="#">Git Tutorial</a> <b>SLIDES</b>	
<b>[Fri] Quantifying Perf 1 (Jan 29)</b> <ul style="list-style-type: none"><li>• check_output</li><li>• time</li></ul> <b>Read:</b> <a href="#">HTML</a> , <a href="#">NB</a> <b>Released:</b> <a href="#">P1 (perf measurements)</a>	

read syllabus carefully  
and checkout other content

I'll also use **Canvas** for four things:

- general announcements
- quizzes
- online office hours
- simple grade summaries (not feedback or exam answers)

# Scheduled Activities

## Lectures

- 3 times weekly
- feel free to bring a laptop
- will generally be recorded+posted online  
(questions will be recorded -- feel free to save until after if you aren't comfortable being recorded)

## Lab

- Weekly on Mondays, bring a laptop
- Work through lab exercises with group mates
- 320 staff will walk around to answer questions
- Required for participation credit!
- Fill "Lab Attendance" (one submission per group) or "Lab Absence" (individual) each week for credit: <https://tyler.caraza-harter.com/cs320/s22/surveys.html>



# Class organization: People

## Teams

- you'll be assigned to a team of 4-7 students
- teams will last the whole semester
- some types of collaboration with team members are allowed (not required) on graded work, such as projects+quizzes
- most collaboration with non-team members is not allowed

## Staff

1. Instructor
2. Teaching Assistants (grad students)
3. Mentors (undergrads)

we all provide office hours, and you can attend any that you prefer!

# Class organization: People

## Teams

- you'll be assigned to a team of 4-7 students
- teams will last the whole semester
- some types of collaboration with team members are allowed (not required) on graded work, such as projects+quizzes
- most collaboration with non-team members is not allowed

## Staff

1. Instructor
2. Teaching Assistants
  - head TA: in charge of projects
  - team TA: primary contact for team, same whole semester
  - grader TA: reviews projects (rotates weekly)
3. Mentors

we all provide office hours, and you can attend any that you prefer!

# Other Communication

## Piazza

- find link on site
- don't post >5 lines of project-related code (considered cheating)

## Forms

- <https://tyler.caraza-harter.com/cs320/s22/surveys.html>
- Who are you? Feedback Form. Thank you! Grading Issues.

## Email

- me: [tharter@wisc.edu](mailto:tharter@wisc.edu)
- TAs: <https://tyler.caraz-harter.com/cs320/s22/contact.html>

# Course Etiquette

## Meetings

1. office hours are drop-in (no need to reserve)
2. email me to schedule individual meetings

## Email

3. let us know your NetID (if not from [netid@wisc.edu](mailto:netid@wisc.edu))
4. don't start new email thread if topic is the same
5. CC team members when appropriate
6. unless urgent, please give me 48 hours to respond before following up (I'll try to be faster usually)
7. use your judgement about whether to email me or TA first  
(if one TA doesn't know something, ask me next before others)
8. if general question, consider using piazza instead if general interest

# Graded Work: Exams/Quizzes

## Ten Online Quizzes - 1% each

- cumulative
- no time limit
- on Canvas, open book/notes
- can take together AT SAME TIME with team members (no other human help allowed)

## Midterms - 10%, 12%

- individual, multi-choice, 40 minutes
- one page notes, both sides
- in class

## Final - 15%

- individual, multi-choice, 2 hours
- one page notes, both sides

# Graded Work: Projects+Participation


## 7 Projects - 7% each

- **format:** notebook, module, or program
- part 1: you can optionally collaborate with team
- part 2: must be individually (only help from 320 staff)
- still a `tester.py`, but more depends on TA evaluation (more plots)
- ask for specific feedback  
(giving constructive criticism is a priority in CS 320)

## Participation - 4%


- lab attendance
- class surveys
- etc.

# Time Commitment



## Observations

- 10-12 hours per project is typical
- 20% of students sometimes spend 20+ hours on some projects
- students who were faster early on were less likely to complete the course



## Weekly Expectations

- 4 hours - lecture/lab
- 6 hours - project coding
- 2 hours - reading/quizzes/etc

# Academic Misconduct

Read syllabus to make sure you know what is and isn't OK.

It's not obvious!

Since Fall 2019, I have made the following misconduct reports:

- **46** students for cheating on projects
- **2** past students for sharing solutions from past semesters
- **7** students for cheating on exams

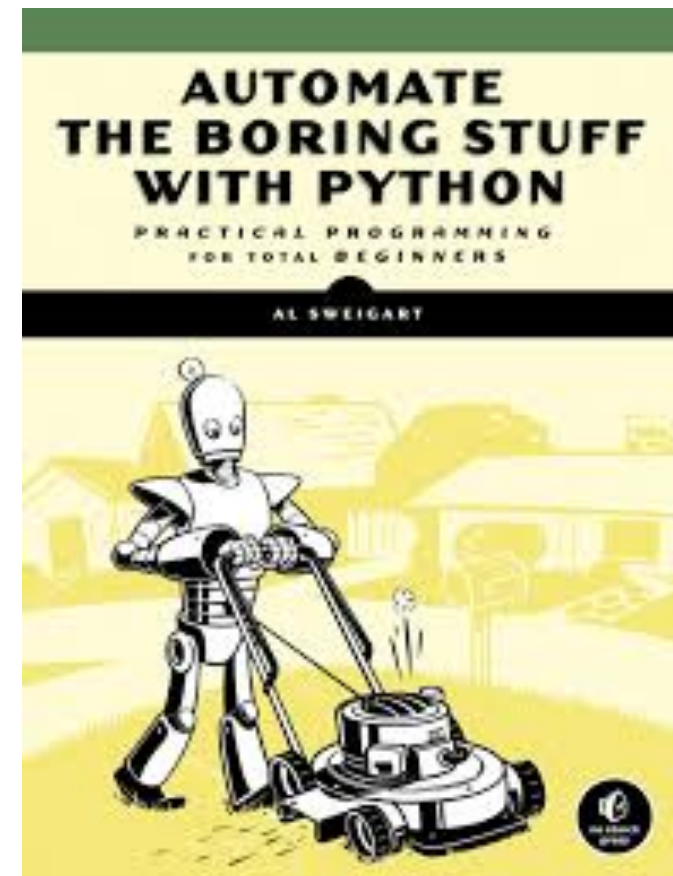
How we'll keep the class fair

- run **MOSS** on submissions
- randomize exam question order

Please talk to me if you're  
feeling overwhelmed with 320  
or your semester in general!



# Reading: same as 220/301 and some others...



I'll post links to other online articles and my own notes

Lectures don't assume any reading prior to class

# Tips for 320 Success

1. Just show up!
  - ➡ Get 100% on participation and don't miss quizzes
2. Use office hours
  - ➡ we're idle after a project release and swamped before a deadline
3. Do labs before projects
4. Take the lead on group collaboration
5. Learn debugging
6. Run the tester often
7. If you're struggling, reach out -- the sooner, the better

Any questions?

# Today's Lecture:

# **Reproducibility**

Reproducibility



 All

 News

 Images

 Books

 Videos

 More

Settings

Tools

About 44,700,000 results (0.64 seconds)

## Dictionary

Search for a word



re·pro·duc·i·bil·i·ty

/ˌrēprəˌd(y)ŏʊəsəˈbɪlədē/

*noun*

noun: **reproducibility**

the ability to be reproduced or copied.

"the reproducibility of reconstructive surgery techniques"

- the extent to which consistent results are obtained when an experiment is repeated.  
"the experiments were conducted numerous times to test the reproducibility of the results"

**Discuss:** *how might we define "reproducibility" for a data scientist?*

**Big question:** *will my program run on someone else's computer?*  
(not necessarily written in Python)

Things to match:

1



Hardware

2

Operating System

3


Dependencies ← next lecture



# Hardware: Mental Model of Process Memory


*Imagine...*

- one huge list, **per each** running program **process**
- every entry in the list is an integer between 0 and 255 (aka a **"byte"**)



How can we use one giant list to handle the following?

- multiple lists
  - variables and other references
  - strings
  - code
- data




*Is this really all we have for state?*



# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- strings
- code



the [3,20] list starts at index address 8 in the giant list

the [11,22,33] list starts at address 12 in the giant list

# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- strings
- code


0	0	0	0	0	0	0	0	3	20	0	0	11	22	33	0	...
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	

*implications for performance...*

```
# fast  
L2.append(44)
```

# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- strings
- code



*implications for performance...*

```
# fast  
L2.append(44)
```

# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- strings
- code

0	0	0	0	0	0	0	0	3	20	0	0	11	22	33	44	...
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	


*implications for performance...*

```
# fast  
L2.append(44)
```

```
# slow  
L2.pop(0)
```

# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- strings
- code




*implications for performance...*

```
# fast  
L2.append(44)
```

```
# slow  
L2.pop(0)
```

# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- strings
- code




We'll think more rigorously about performance in CS 320 (big-O notation)

```
# fast  
L2.append(44)
```

```
# slow  
L2.pop(0)
```

# How can we use one giant list to handle the following?


- multiple lists
- **variables and other references**
- strings
- code



Python 3.6

```
1 x = [3, 20]
2 y = x
```

[Edit this code](#)




Python Tutor's visualization

How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- **strings**
- code

discuss: how?




*Is this really all we have for state?*



# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- **strings**
- code




encoding:

code	letter
65	A
66	B
67	C
68	D
...	...

```
f = open("file.txt", encoding="utf-8")
```

# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- **strings**
- code



encoding:


code	letter
65	A
66	B
67	C
68	D
...	...

```
f = open("file.txt", encoding="utf-8")
```

# How can we use one giant list to handle the following?

- multiple lists
- variables and other references
- strings
- **code**

```
while ????:  
    i += 2  
    # what line next?
```





Instruction Set	code	operation
	5	ADD
	8	SUB
	33	JUMP
	...	...

# Hardware: Mental Model of CPU

CPUs interact with memory:

- keep track of what instruction we're on
- understand instruction codes
- much more



Write code in Python 3.6

(drag lower right corner to resize code editor)

```
→ 1 XXXXXXXXXX
  2 XXXXXXXXXX
  3 XXXXXXXXXX
```

→ line that just executed

→ next line to execute


Instruction Set

code	operation
5	ADD
8	SUB
33	JUMP
...	...


# Hardware: Mental Model of CPU

CPU's interact with memory:

- keep track of what instruction we're on
- understand instruction codes
- much more



add 2 to variable




Instruction Set	code	operation
	5	ADD
	8	SUB
	33	JUMP
	...	...


# Hardware: Mental Model of CPU

CPU's interact with memory:

- keep track of what instruction we're on
- understand instruction codes
- much more



go back to top of loop


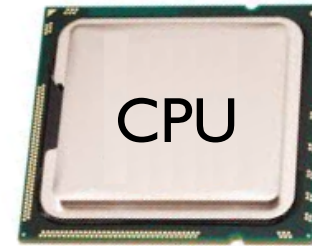


Instruction Set	code	operation
	5	ADD
	8	SUB
	33	JUMP
	...	...

# Hardware: Mental Model of CPU

CPU's interact with memory:

- keep track of what instruction we're on
- understand instruction codes
- much more



Instruction Set	code	operation
	5	ADD
	8	SUB
	33	JUMP
	...	...

# Hardware: Mental Model of CPU

discuss: what would happen if a  
CPU tried to execute an  
instruction for a different CPU?

0	0	0	0	0	0	0	0	...	...	5	2	33	8	0	0	...
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	

Instruction Set  
for CPU X

code	operation
5	ADD
8	SUB
33	JUMP
...	...



Instruction Set  
for CPU Y

code	operation
5	SUB
8	ADD
33	undefined
...	...



# Hardware: Mental Model of CPU

a CPU can only run programs that use instructions it understands!







Instruction Set  
for **CPU X**

code	operation
5	ADD
8	SUB
33	JUMP
...	...

Instruction Set  
for **CPU Y**

code	operation
5	SUB
8	ADD
33	undefined
...	...

# A Program and CPU need to "fit"




# A Program and CPU need to "fit"



*why haven't we noticed this yet  
for our Python programs?*


# Interpreters



**Interpreters** (such as `python.exe`) make it easier to run the same code on different machines

A **compiler** is another tool for running the same code on different CPUs

# Interpreters



**Interpreters** (such as python.exe) make it easier to run the same code on different machines

**Discuss:** *if all CPUs had the instruction set, would we still need a Python interpreter?*

**Big question:** *will my program run on someone else's computer?*  
(not necessarily written in Python)

Things to match:

1



Hardware

2

Operating System

3

Dependencies ← next lecture



**Big question:** *will my program run on someone else's computer?*  
(not necessarily written in Python)

Things to match:

1


Hardware

2

Operating System

3

Dependencies ← next lecture



Linux™

many others...



**Red Hat**



ANDROID



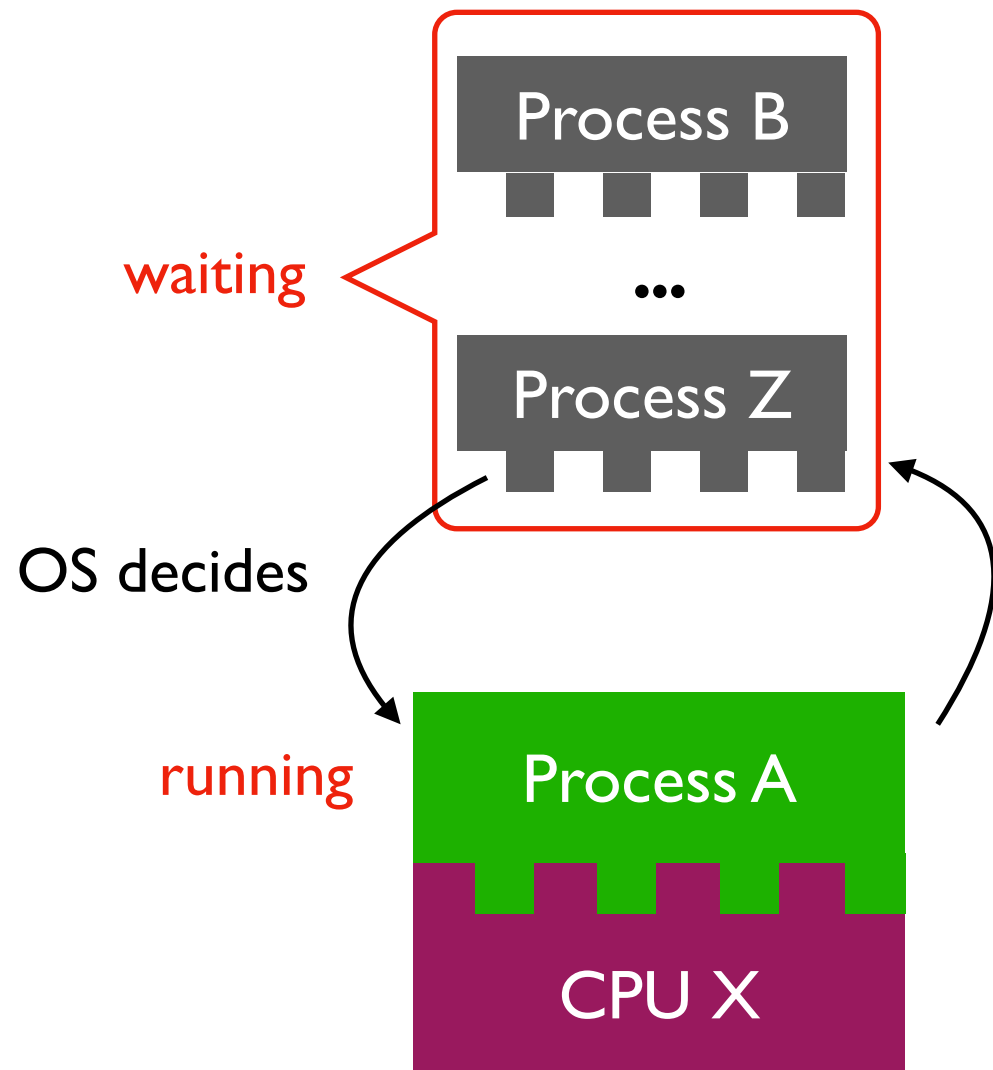
ubuntu.

[this semester]

# OS jobs: Allocate and Abstract Resources

[like CPU, hard drive, etc]

## 1 Allocation



only one process can run on CPU at a time  
(or a few things if the CPU has multiple "cores")

## 2 Abstraction

```
f = open("file.txt")  
data = f.read()  
f.close()
```

convenient

Operating System


inconvenient



ignorant of  
files/directories



# Harder to reproduce on different OS...




```
f = open("/data/file.txt")  
...
```

The Python interpreter mostly lets you  
[Python Programmer] ignore the CPU you run on.

But you still need to work a bit to "fit" the code to the OS.

# Harder to reproduce on different OS...




```
f = open("c:\data\file.txt")  
...
```

The Python interpreter mostly lets you [Python Programmer] ignore the CPU you run on.

But you still need to work a bit to "fit" the code to the OS.

# Harder to reproduce on different OS...



# solution 1:

```
f = open(os.path.join("data", "file.txt"))  
...
```

# solution 2:

tell anybody reproducing your results to use the same OS!


tradeoffs?

The Python interpreter mostly lets you  
[Python Programmer] ignore the CPU you run on.

But you still need to work a bit to "fit" the code to the OS.

# VMs (Virtual Machines)

popular virtual  
machine software



With the right virtual machines created and operating systems installed, you could run programs for Mac, Linux, and Windows -- at the same time without rebooting!

# The Cloud


cloud providers let you rent VMs  
in the cloud on hourly basis  
(e.g., \$15 / month)

popular cloud providers



Google Cloud Platform

we'll use GCP virtual  
machines this semester  
[setup in lab]



← Windows, Mac, whatever

`ssh user@best-linux.cs.wisc.edu` ← run in PowerShell/  
bash to access CS lab

# Lecture Recap: Reproducibility

**Big question:** *will my program run on someone else's computer?*

Things to match:

1

Hardware

← a program must fit the CPU;  
`python.exe` will do this, so  
`program.py` won't have to

2

Operating System

← we'll use Ubuntu Linux on  
virtual machines in the cloud

3

Dependencies

← next time: versioning

# Recap of 15 new terms

**reproducibility:** others can run our analysis code and get same results

**process:** a running program

**byte:** integer between 0 and 255

**address space:** a big "list" of bytes, per process, for all state

**address:** index in the big list

**encoding:** pairing of ~~letters~~ characters with numeric codes

**CPU:** chip that executes instructions, tracks position in code

**instruction set:** pairing of CPU instructions/ops with numeric codes

**operating system:** software that allocates+abstracts resources

**resource:** time on CPU, space in memory, space on SSD, etc

**allocation:** the giving of a resource to a process

**abstraction:** hiding inconvenient details with something easier to use

**virtual machine:** "fake" machine running on ~~real~~ physical machine

allows us to run additional operating systems

**cloud:** place where you can rent virtual machines and other services

**ssh:** secure shell -- tool that lets you remotely access another machine